

# Hive Practicals

## Quick setup (session & DB)

```
-- start a session in beeline or hive CLI;  
create/use DB  
CREATE DATABASE IF NOT EXISTS demo_hive;  
USE demo_hive;
```

## 1. Create managed table (TEXTFILE)

```
CREATE TABLE employees(  
    id INT,  
    name STRING,  
    dept STRING,  
    salary DOUBLE,  
    doj STRING  
)  
ROW FORMAT DELIMITED  
FIELDS TERMINATED BY ','  
STORED AS TEXTFILE;
```

## 2. Create external table

```
CREATE EXTERNAL TABLE ext_events(  
    event_id STRING,
```

```
    ts STRING,  
    payload STRING  
)  
ROW FORMAT DELIMITED FIELDS TERMINATED BY '\t'  
LOCATION '/user/hive/external/events';
```

### 3. Create table as select (CTAS)

```
CREATE TABLE top_paid AS  
SELECT id, name, salary FROM employees WHERE salary  
> 100000;
```

### 4. Load data (local -> Hive)

```
-- loads local file into managed table  
LOAD DATA LOCAL INPATH '/local/path/employees.csv'  
OVERWRITE INTO TABLE employees;  
-- or load HDFS file into external table path (no  
LOAD needed if file in LOCATION)
```

### 5. Partitions (static)

```
CREATE TABLE sales(  
    order_id STRING, amount DOUBLE  
)  
PARTITIONED BY (year INT, month INT)  
STORED AS PARQUET;  
  
-- add partition and load  
ALTER TABLE sales ADD PARTITION (year=2024, month=8)  
LOCATION '/data/sales/2024/08';
```

## 6. Dynamic partition insert

```
SET hive.exec.dynamic.partition=true;
SET hive.exec.dynamic.partition.mode=nonstrict;

INSERT INTO TABLE sales PARTITION (year, month)
SELECT order_id, amount, year(order_dt),
month(order_dt) FROM staging_orders;
```

## 7. Bucketing

```
CREATE TABLE users_bucketed(
    user_id INT,
    name STRING
)
CLUSTERED BY (user_id) INTO 8 BUCKETS
STORED AS ORC;

-- enable bucketing reads/writes
SET hive.enforce.bucketing=true;
INSERT INTO TABLE users_bucketed SELECT * FROM
users;
```

## 8. ORC / Parquet / SequenceFile example

```
CREATE TABLE orc_table(id INT, data STRING) STORED
AS ORC;
```

```
CREATE TABLE parquet_table(id INT, data STRING)
STORED AS PARQUET;
CREATE TABLE seq_table(id INT, data STRING) STORED
AS SEQUENCEFILE;
```

## 9. Compressed ORC/Parquet writes

```
SET parquet.compression=SNAPPY;
SET hive.exec.compress.output=true;
SET hive.exec.orc.compression=ZLIB;
```

```
INSERT OVERWRITE TABLE orc_table SELECT * FROM
some_source;
```

## 10. Partition pruning demonstration

```
-- assume sales is partitioned by year,month
SELECT SUM(amount) FROM sales WHERE year=2024 AND
month=8; -- prunes partitions
EXPLAIN SELECT SUM(amount) FROM sales WHERE
year=2024 AND month=8;
```

## 11. Basic SELECT, WHERE, GROUP BY, HAVING

```
SELECT dept, COUNT(*) cnt, AVG(salary) avg_sal
FROM employees
WHERE salary > 20000
GROUP BY dept
HAVING COUNT(*) > 2
```

```
ORDER BY avg_sal DESC;
```

## 12. Joins (inner, left, right, full)

```
SELECT e.id, e.name, d.dept_name  
FROM employees e  
JOIN departments d ON e.dept = d.dept_code;
```

```
-- left join  
SELECT * FROM employees e LEFT JOIN dept_budget b ON  
e.dept=b.dept_code;
```

## 13. Map-side join (mapjoin hint) for small table

```
-- force the smaller table into memory to avoid  
shuffle  
SELECT /*+ MAPJOIN(departments) */ e.id, e.name,  
d.dept_name  
FROM employees e JOIN departments d ON e.dept =  
d.dept_code;
```

## 14. Lateral view + explode (UDTF use)

```
CREATE TABLE user_tags(user_id INT, tags  
ARRAY<STRING>) STORED AS ORC;
```

```
-- explode tags to rows  
SELECT u.user_id, t.tag  
FROM user_tags u
```

```
LATERAL VIEW EXPLODE(tags) exploded_table as tag;
```

## 15. Window functions

```
SELECT id, name, dept, salary,  
       ROW_NUMBER() OVER (PARTITION BY dept ORDER BY  
salary DESC) rn,  
       RANK() OVER (PARTITION BY dept ORDER BY  
salary DESC) rnk,  
       NTILE(4) OVER (ORDER BY salary) quartile  
FROM employees;
```

## 16. Grouping sets / rollup / cube

```
-- rollup  
SELECT dept, year, SUM(amount)  
FROM sales  
GROUP BY ROLLUP(dept, year);
```

```
-- cube  
SELECT dept, year, SUM(amount)  
FROM sales  
GROUP BY CUBE(dept, year);
```

```
-- grouping sets  
SELECT dept, year, SUM(amount)  
FROM sales  
GROUP BY GROUPING SETS ((dept, year), (dept),  
(year), ());
```

## 17. Subqueries (IN / EXISTS / scalar)

```
SELECT name FROM employees WHERE id IN (SELECT  
emp_id FROM managers);
```

```
SELECT * FROM departments d WHERE EXISTS (SELECT 1  
FROM employees e WHERE e.dept=d.dept_code);
```

## 18. Views and materialized views

```
CREATE VIEW v_high_salary AS SELECT id,name,salary  
FROM employees WHERE salary>100000;
```

```
CREATE MATERIALIZED VIEW mv_sales_monthly  
AS SELECT year,month, SUM(amount) total FROM sales  
GROUP BY year,month;  
-- refresh depends on Hive version: REFRESH  
MATERIALIZED VIEW mv_sales_monthly;
```

## 19. Transactions and ACID (insert/update/delete)

```
-- requires transactional table and ACID settings:  
SET hive.support.concurrency=true;  
SET hive.enforce.bucketing=true;  
SET hive.exec.dynamic.partition.mode=nonstrict;  
SET  
hive.txn.manager=org.apache.hadoop.hive.q1.lockmgr.D  
bTxnManager;  
SET hive.compactor.initiator.on=true;  
SET hive.compactor.worker.threads=1;
```

```
CREATE TABLE txn_users(  
    id INT,  
    name STRING  
) CLUSTERED BY (id) INTO 4 BUCKETS  
STORED AS ORC  
TBLPROPERTIES ('transactional'='true');  
  
-- DML  
INSERT INTO txn_users VALUES (1,'Alice');  
UPDATE txn_users SET name='Alicia' WHERE id=1;  
DELETE FROM txn_users WHERE id=1;
```

## 20. INSERT OVERWRITE vs INSERT INTO

```
-- overwrite entire table or partition  
INSERT OVERWRITE TABLE sales PARTITION (year=2024,  
month=8)  
SELECT * FROM staging_sales WHERE year=2024 AND  
month=8;  
  
-- append  
INSERT INTO TABLE sales PARTITION (year, month)  
SELECT order_id, amount, year(order_dt),  
month(order_dt) FROM staging_orders;
```

## 21. UDF (Java/JS) — example registering and using a simple UDF



## (assumes udf jar already on HDFS/local)

```
-- add jar (local or hdfs)
ADD JAR /path/to/myudfs.jar;

-- create temporary function
CREATE TEMPORARY FUNCTION reverse_str AS
'com.example.hive.udf.ReverseStringUDF';

-- use
SELECT reverse_str(name) FROM employees LIMIT 5;
```

(If you want, I can provide sample Java code for a UDF/UDAF/UDTF.)

## 22. UDAF / UDTF brief examples

```
-- UDAF usage: assume a jar provides 'median_udaf'
ADD JAR /path/median_udaf.jar;
CREATE TEMPORARY FUNCTION median AS
'com.example.hive.udaf.Median';
SELECT dept, median(salary) FROM employees GROUP BY
dept;

-- UDTF example: explode() is a built-in UDTF used
with LATERAL VIEW above
```

## 23. SerDe usage (custom delimited or JSON SerDe)

```
-- JSON SerDe (builtin or external) example:
```

```
CREATE TABLE events_json(  
    event_id STRING,  
    ts STRING,  
    payload MAP<STRING,STRING>  
)  
ROW FORMAT SERDE  
'org.apache.hive.hcatalog.data.JsonSerDe'  
STORED AS TEXTFILE;
```

## 24. ANALYZE / COMPUTE STATS & optimizer

```
ANALYZE TABLE employees COMPUTE STATISTICS;  
ANALYZE TABLE sales PARTITION(year,month) COMPUTE  
STATISTICS;  
-- helps cost-based optimizer and partition pruning
```

## 25. EXPLAIN & PROFILE queries

```
EXPLAIN FORMATTED SELECT dept, AVG(salary) FROM  
employees GROUP BY dept;  
-- or use EXPLAIN EXTENDED / EXPLAIN COST if  
supported
```

## 26. EXTERNAL TABLE over HDFS/Parquet/ORC files (schema-on-read)

```
CREATE EXTERNAL TABLE parquet_ext(  

```

```
id INT, name STRING  
) STORED AS PARQUET LOCATION '/data/parquet/users';
```

## 27. File format-specific operations (ORC stats, vectorized reads)

```
SET hive.vectorized.execution.enabled=true;  
SET hive.vectorized.execution.reduce.enabled=true;  
-- ORC statistics example when creating ORC to get  
min/max stats  
INSERT OVERWRITE TABLE orc_table SELECT * FROM  
employees;
```

## 28. Sampling

```
SELECT * FROM employees TABLESAMPLE(10 PERCENT);  
-- or: SELECT * FROM employees TABLESAMPLE(BUCKET 3  
OUT OF 32);
```

## 29. UNION and UNION ALL

```
SELECT id,name FROM employees WHERE dept='HR'  
UNION ALL  
SELECT id,name FROM contractors WHERE dept='HR';
```

## 30. SHOW / DESCRIBE / MSCK (repair) / ALTER

```
SHOW TABLES;  
DESCRIBE EXTENDED employees;  
ALTER TABLE sales ADD PARTITION(year=2023,  
month=12);  
MSCK REPAIR TABLE ext_events; -- recovers partitions  
from HDFS for external tables
```

## **31. Indexes (historical / deprecated; still available in some versions)**

```
-- note: indexes are rarely used and often  
deprecated; example:  
CREATE INDEX idx_emp_dept ON TABLE employees(dept)  
AS 'COMPACT' WITH DEFERRED REBUILD;  
ALTER INDEX idx_emp_dept ON employees REBUILD;
```

## **32. Hive configuration settings / session tuning**

```
-- useful settings  
SET hive.exec.reducers.bytes.per.reducer=268435456;  
-- 256MB  
SET hive.auto.convert.join=true; -- convert to  
mapjoin if small table  
SET hive.exec.parallel=true;
```

## **33. Using BEELINE (JDBC) example**

```
beeline -u "jdbc:hive2://namenode:10000/default" -n  
hiveuser -p hivepassword
```

## 34. Export / Import table data

```
EXPORT TABLE employees TO  
'/backup/employees_export';  
IMPORT TABLE employees FROM  
'/backup/employees_export';
```

## 35. Encryption / ACL hints (high-level)

- Hive itself integrates with HDFS encryption zones and external security (Kerberos, Ranger, Sentry). Typical steps: enable Kerberos, configure Ranger policies, ensure metastore/hive-server2 run under Kerberos principals. (Exact commands depend on your cluster/security tools.)